

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE: serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE: impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



Synchronization

Barriers

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

License Information

This work is licensed under the license



Attribution-NonCommercial-NoDerivatives 4.0 International

This license requires that reusers give credit to the creator. It allows reusers to copy and distribute the material in any medium or format in unadapted form and for noncommercial purposes only.

① **BY:** Credit must be given to you, the creator.

② **NC:** Only noncommercial use of your work is permitted.

Noncommercial means not primarily intended for or directed towards commercial advantage or monetary compensation.


③ **ND:** No derivatives or adaptations of your work are permitted.

To view a copy of the license, visit:

<https://creativecommons.org/licenses/by-nc-nd/4.0/?ref=chooser-v1>

Barriers

- ❖ Barriers can be used to coordinate multiple threads working in parallel
 - A barrier allows each thread to wait until all cooperating threads have reached the same point, and then continue executing from there



Barriers

- ❖ Barriers generalize the `t.join()` function
 - The function **`t.join()`** acts as a barrier to allow **one** thread to wait until another thread completes its processing
 - Barriers allow **an arbitrary number** of threads to wait until all of the threads have completed their processing
 - The threads don't have to exit, as they can continue working after all threads have reached the barrier

Trivial solution

❖ A possible trivial solution


- Use one semaphore for each thread T_i
- Implement one for the extra process B (for the barrier) using one more semaphore

It uses too many semaphores

The barrier process waits n times on its semaphores and then wakes-up all threads T_i

B (pseudo-code)

```
for (i=0; i<n; i++)  
    wait (semB);  
for (i=0; i<n; i++)  
    signal (sem[i]);
```



T_i (pseudo-code)

```
signal (semB);  
wait (sem[i]);  
...
```

Each T_i wakes-up the barrier process and wait on its own semaphore

Barriers in C++

- ❖ In C++, synchronization barriers can be implemented using several primitives, with notable additions in C++20
 - A `std::latch` is a one-time synchronization primitive
 - It counts down from an initial value set in its constructor
 - Ideal when threads need to wait for a specific number of tasks to complete before proceeding

Methods	Meaning
<code>wait()</code>	To wait at the barrier until the counter reaches zero.
<code>count_down()</code>	Decrement the counter without waiting.

Barriers in C++

- A `std::barrier` is similar to `std::latch`, but **reusable**
 - After the counter reaches zero, it resets to its initial value, allowing for repeated synchronization
 - Suitable for phased computations where threads need to synchronize multiple times

Methods	Meaning
<code>arrive_and_wait()</code>	To wait at the barrier.
<code>arrive_and_drop()</code>	To remove a thread from the barrier without waiting.

Barriers in C++

- A `std::flex_barrier` is an extension of `std::barrier`, allowing **the execution of a function** when the counter reaches zero
 - This function can be specified during barrier construction.

Example

Main program calling
the threads

```
#include <iostream>
#include <barrier>
#include <thread>
#include <vector>
```

Used for a general
I/O critical section

```
std::mutex mtx_out;
void worker(std::barrier<> &);
```

Define
the barrier for
10 threads

```
int main() {
    const int numThreads = 10;
    std::barrier<> barrier(numThreads);
    std::vector<std::thread> threads;
    for (int i = 0; i < numThreads; ++i) {
        threads.emplace_back(worker, std::ref(barrier));
    }
    for (auto &thread : threads) {
        thread.join();
    }
    return 0;
}
```

Example

Use the barrier to synch
numThreads **once**

```
void worker(std::barrier<> &barrier) {  
    // Perform some work  
    mtx_out.lock();  
    std::cout << "Thread performing work...\n";  
    mtx_out.unlock();  
  
    // Wait for all threads to reach this point  
    barrier.arrive_and_wait();  
  
    // Continue after all threads have synchronized  
    mtx_out.lock();  
    std::cout << "Thread continuing after barrier.\n";  
    mtx_out.unlock();  
  
    return;  
}
```

The mutex
protect cout

The mutex
protect cout

Example

Use the barrier to synch
iterations **times**

```
void worker(std::barrier<> &barrier, int id,
            int iterations) {
    for (int i = 0; i < iterations; ++i) {
        mtx_out.lock();
        cout << "Thread " << id << " it " << i << endl;
        mtx_out.unlock();

        barrier.arrive_and_wait();

        mtx_out.lock();
        cout << "Thread " << id <<
            " after barrier on " << i << endl;
        mtx_out.unlock();
    }
    return;
}
```

The barrier does not
have to be re-initialized

The mutex
protect cout

Conclusions

- ❖ Barriers are used to coordinate multiple threads working in parallel
 - You want all threads to wait until everyone has arrived at a certain point
 - A simple semaphore would do the exact opposite, i.e., each thread would keep running and the last one will go to sleep

Exercise 01

Acyclic barrier


- ❖ Suppose barrier constructs do not exist
 - Re-implement them using only **one** semaphore **and** one mutex

```
std::barrier<> barrier(numThreads);  
std::vector<std::thread> threads;  
for (int i = 0; i < numThreads; ++i) {  
    threads.emplace_back(worker, std::ref(barrier));  
}  
for (auto &thread : threads) {  
    thread.join();  
}
```

```
void worker(std::barrier<> &barrier) {  
    ...  
    barrier.arrive_and_wait();  
    ...  
}
```

Threads
(**acyclic** behavior)

Synchronization point
among all threads



Solution 01

```
#include <mutex>
#include <semaphore>
#include <thread>
#include <vector>
#include <iostream>

const int numThreads = 10;
std::mutex mtx;
std::mutex mtx_out;
std::counting_semaphore<numThreads> sem{0};
int count;    // Total number of threads
int arrived; // Number of threads that have arrived

void worker(int);
```


Solution 01

```
int main() {  
    count = numThreads;  
    arrived = 0;  
  
    std::vector<std::thread> threads;  
    for (int i = 0; i < numThreads; ++i) {  
        threads.emplace_back(worker, i);  
    }  
  
    for (auto &thread : threads) {  
        thread.join();  
    }  
  
    return 0;  
}
```

Solution 01

```
void worker(int id) {  
    mtx_out.lock();  
    std::cout << "Thread " << id << " work..." << std::endl;  
    mtx_out.unlock();  
    mtx.lock();  
    arrived++;  
    if (arrived == count) {  
        for (int i=0; i<count; ++i) sem.release();  
    }  
    mtx.unlock();  
    sem.acquire();  
    mtx_out.lock();  
    std::cout << "Thread " << id << " after barrier." << std::endl;  
    mtx_out.unlock();  
    return;  
}
```

Protect counter

Un-protect counter

Solution 01

Solution with turnstile

```
void worker(int id) {  
    mtx_out.lock();  
    std::cout << "Thread " << id << " work..." << std::endl;  
    mtx_out.unlock();  
    mtx.lock();  
    arrived++;  
    if (arrived == count)  
        sem.release();  
    mtx.unlock();  
    sem.acquire();  
    sem.release();  
    mtx_out.lock();  
    std::cout << "Thread " << id << " after barrier." << std::endl;  
    mtx_out.unlock();  
    return;  
}
```

Protect counter

Un-protect counter

Turnstile

One **extra** sem_post is
done (pay attention to
cycling threads)

Exercise 02

Cyclic barrier


- ❖ Suppose barrier constructs do not exist
 - Re-implement them using only **one** semaphore **and** one mutex

```
std::barrier<> barrier(numThreads);  
std::vector<std::thread> threads;  
for (int i = 0; i < numThreads; ++i) {  
    threads.emplace_back(worker, std::ref(barrier));  
}  
for (auto &thread : threads) {  
    thread.join();  
}
```

```
void worker(std::barrier<> &barrier) {  
    for (int i = 0; i < iterations; ++i) {  
        ...  
        barrier.arrive_and_wait();  
        ...  
    }  
}
```

Threads
(cyclic behavior)

Synchronization point
among all threads



Buggy solution

Buggy Solution

```
void worker(int id) {
    for (int i=0; i<iterations; i++) {
        mtx_out.lock();
        std::cout << "Thread " << id << " work..." << std::endl;
        mtx_out.unlock();
        mtx.lock();
        arrived++;
        if (arrived == count) {
            arrived = 0;
            for (int i=0; i<count; ++i) sem.release();
        }
        mtx.unlock();
        sem.acquire();
        mtx_out.lock();
        std::cout << "Thread " << id << " continuing after barrier." << std::endl;
        mtx_out.unlock();
    }
    return;
}
```

Last threads awakes all

Waiting point for all threads

A fast threads can cycle more than once !

Solution

Correct solution
(part I)

```
void worker(int id) {  
    for (int i=0; i<iterations; i++) {  
        mtx_out.lock();  
        std::cout << "Thread " << id <<  
            " performing work..." << std::endl;  
        mtx_out.unlock();  
        mtx.lock();  
        arrived++;  
        if (arrived == count) {  
            for (int i=0; i<count; ++i) sem1.release();  
        }  
        mtx.unlock();  
        sem1.acquire();  
        mtx_out.lock();  
        std::cout << "Thread " << id <<  
            " moving from B1 to B2." << std::endl;  
        mtx_out.unlock();  
    }  
}
```


Barrier #1

Correct solution
(part II)

```
mtx.lock();
arrived--;
if (arrived == 0) {
    for (int i=0; i<count; ++i) sem2.release();
}
mtx.unlock();
sem2.acquire();
mtx_out.lock();
std::cout << "Thread " << id <<
    " continuing after barrier." << std::endl;
mtx_out.unlock();
}

return;
}
```

Barrier #2



Exercise 03


Concurrent Bubble-sort

- ❖ Write a version of the exchange (bubblesort) sorting algorithm) as follows
 - A static array include n integer values
 - We want to sort it using n identical threads
 - Each thread is in charge of sorting two adjacent elements
 - Thread 0 sort elements 0 and 1
 - Thread 1 sort elements 1 and 2
 - ...
 - Thread $n-1$ sort elements $n-1$ and n

Exercise 03

➤ Each thread

- Compare the two elements it deals with, and exchange them if they are not in the correct order
- Once their work is finished, all the threads wait for each-other, and if
 - All the elements are correctly ordered, the program terminates
 - Otherwise, all threads are run again to make a new series of exchanges




Exercise 03

➤ Each thread

- Compare the two elements it deals with, and exchange them if they are not in the correct order
- Once their work is finished, all the threads wait for each-other, and if
 - All the elements are correctly ordered, the program terminates
 - Otherwise, all threads are run again to make a new series of exchanges

As the order in which all swaps are performed is not defined (inner iteration) the number of necessary outer iterations is upper bounded by n

```
for (i=0; i<n-1; i++)  
    for (j=0; j<n-i-1; j++)  
        if (v[j] > v[j+1])  
            swap (v, i, j+1);
```




Solution

Main (Estract) Part I

```
#include <thread>
#include <vector>
#include <iostream>
#include <barrier>
#include <mutex>

void sorter (int *vet, int vet_size, int i,
             std::vector<std::mutex> &me, std::barrier<> &barrier,
             std::mutex &mtx_out, bool &sorted, bool &all_ok);

int main(int argc, char** argv) {
    if (argc != 2) {
        std::cout << "Syntax: " << argv[0] << " vet_size\n";
        return 1;
    }
}
```



Solution

Main (Extract) Part II

```
const int vet_size = std::stoi(argv[1]);
const int num_threads = vet_size - 1;
std::mutex mtx_out;
std::barrier<> barrier(num_threads);
std::vector<std::mutex> me(vet_size);
int *vet = new int[vet_size];
bool sorted = false;
bool all_ok = true;


std::srand(std::time(nullptr));

for (int i=0; i<vet_size; i++) {
    vet[i] = std::rand() % 1000;
    std::cout << vet[i] << " ";
}

std::cout << "\n";
```

Create a mutex for each element of the vector

Fill the vector with random numbers




Solution

Main (Extract)
Part III

```
std::vector<std::thread> threads;  
for (int i=0; i<num_threads; i++) {  
    threads.emplace_back(sorter, vet, vet_size, i,  
        std::ref(me), std::ref(barrier),  
        std::ref(mtx_out), std::ref(sorted), std::ref(all_ok));  
}  
  
for (auto& thread : threads) {  
    thread.join();  
}  
  
delete[] vet;  
  
return 0;  
}
```

Run threads

Wait for threads




Solution

Worker
Part 1

```
void sorter (int *vet, int vet_size, int i,  
             std::vector<std::mutex> &me, std::barrier<> &barrier,  
             std::mutex &mtx_out, bool &sorted, bool &all_ok) {  
    while (!sorted) {  
        {  
            std::lock_guard<std::mutex> lock1 (me[i]);  
            std::lock_guard<std::mutex> lock2 (me[i+1]);  
            if (vet[i] > vet[i+1]) {  
                mtx_out.lock();  
                mtx_out.unlock();  
                int tmp = vet[i];  
                vet[i] = vet[i + 1];  
                vet[i + 1] = tmp;  
                all_ok = false;  
            }  
        }  
    }  
}
```

Acquire mutexes
for two contiguous
elements

Swap elements
if necessary



Solution

Worker Part II

```
barrier.arrive_and_wait();  
if (i==0) {  
    for (int i=0; i<vet_size; i++) {  
        std::cout << vet[i] << " ";  
    }  
    std::cout << std::endl;  
    if (all_ok)  
        sorted = true;  
    else  
        all_ok = true;  
}  
barrier.arrive_and_wait();  
}  
return;  
}
```

If no exchanges have been done, stop the process

